

First steps in MPI

Joachim Hein (LUNARC, Lund University)

Xin Li (PDC, KTH)

Viktor Rehnberg (NAISS & Chalmers)

Overview

- This lecture: Basic set up inside the code for MPI
- Header files
- Initialisation of the MPI library
- Finalisation of the MPI library

Header files

- Every compilation module accessing MPI requires inclusion of a header file:

- F77 style:

```
include "mpif.h"
```

- Fortran90:

```
use mpi
```

- Fortran08, **MPI 3.0**, Fortran standard compliant!

```
use mpi_f08
```

- C:

```
#include "mpi.h"
```

Importing mpi4py in Python

- Python code does not require compilation or header file, but the MPI module needs to be imported

- Python:

```
from mpi4py import MPI
```

MPI command in C

- In C all MPI commands are functions with return type **int**

```
int MPI_Abdef( arguments )
```

- The returned value is the error code
 - Detailing problems with the command
- Typically very hard to recover from MPI-errors
- Most codes do not check these error codes
- **Rem:** MPI commands can modify arguments
 - pass a pointer

MPI command in Fortran

- In Fortran all MPI commands are subroutines

`MPI_ABCDEF(arguments, ierror)`

- MPI commands in Fortran carry **one more argument** than their C counter part
 - This is optional in Fortran 2008
 - This is of type `int` and returns the error code
- Again, this is typically unchecked, hence easily forgotten while coding
- Forgetting this in F77/F90 typically leads to segmentation faults at runtime

MPI command in Python

- In Python all MPI commands are methods of an MPI communicator

`comm.method(arguments)`

- In general fewer arguments are needed, compared to C and Fortran
- Communication of generic Python objects is done via all-lowercase methods (e.g. `comm.send(...)`)
- Communication of buffer-like objects is done via methods with an uppercase letter (e.g. `comm.Send(...)`)

C++ bindings: depreciated/removed

- MPI used to have special C++ bindings
- Depreciated since MPI standard 2.2 September 2009
- Removed in MPI standard 3.0 September 2012
- Use C bindings in C++ programs
 - Consider wrapping in OO-style for your app's needs

MPI_Init

- The first MPI call of any MPI program has to be **MPI_Init**
- In C:

```
int MPI_Init(int *argc, char ***argv)
```

- Arguments are same as **main**
- Alternatively modern MPI libraries allow to pass **Null**

- In Fortran

```
MPI_INIT(IERROR)
```

```
INTEGER :: IERROR
```

MPI_Finalize

- The last MPI call has to be **MPI_Finalize**
- In C:

```
int MPI_Finalize(void)
```

- In Fortran:

```
MPI_FINALIZE(IERROR)  
INTEGER :: IERROR
```

No init and finalize required in Python

- In Python, MPI is initialized upon import, and finalized upon exit

```
from mpi4py import MPI
```

Minimal program in C

```
#include "mpi.h"

int main(int argc, char **argv)
{
    MPI_Init(&argc, &argv); // alt.: NULL,NULL

    // further MPI calls go here!

    MPI_Finalize();
}
```

Minimal program in Fortran

```
program minimpi
  use mpi      ! alt.: include "mpif.h"
  integer:: ierror

  call MPI_INIT(ierror)
              ! further MPI calls go here
  call MPI_FINALIZE(ierror)

end program minimpi
```

Minimal program in Python

```
from mpi4py import MPI
```

```
# further MPI calls go here
```

Summary

- Basic requirements for an MPI program
 - Header files
 - Initialising MPI
 - Finalising MPI

Communicators

Overview


- Concept of communicators
- Predefined communicator
- Querying basic properties of the communicator

Communicator

- Most messages passed inside (intra-)communicator
- Communicator
 - Group of processes
 - Stores communication universe
 - Order
 - Can have additional topology


Example communicator

- Picture shows:
 - 9 processes
 - 3 communicators
- Processes carry label
 - Here: labels for red communicator
- Labels start at 0




Example communicator

- Picture shows:
 - 9 processes
 - 3 communicators
- Processes carry label
 - Here: labels for violet communicator
- Labels start at 0



Example communicator

- Picture shows:
 - 9 processes
 - 3 communicators
- Processes carry label
 - Here: labels for blue communicator
- Labels start at 0
- **Label depends on the communicator**



Predefined communicator


MPI_COMM_WORLD

- After **MPI_Init** one predefined communicator:

MPI_COMM_WORLD

(Python: **MPI.COMM_WORLD**)

- This contains all processes
 - In C, this is typedef: **MPI_Comm**
 - In Fortran90 this is: **INTEGER**
 - In Fortran 2008 this is: **type (MPI_Comm)**
- Further communicators: user created



MPI_Comm_size

- Number of processes in a communicator
- In C:

```
int MPI_Comm_size(MPI_Comm comm, int *size)
```

- In Fortran 90:

```
MPI_COMM_SIZE(COMM, SIZE, IERROR)  
INTEGER :: COMM, SIZE, IERROR)
```

- Arguments:
 - comm**: communicator (input)
 - size**: number of processes (output)

Get_size in Python

- Number of processes in a communicator

```
comm.Get_size()
```

- No arguments
- Returns the number of processes in a communicator
- (alt. use `comm.size`)

MPI_Comm_rank

- Rank (label) of the process
- In C:

```
int MPI_Comm_rank(MPI_Comm comm, int *rank)
```

- In Fortran 90:

```
MPI_COMM_RANK(COMM, RANK, IERROR)  
INTEGER :: COMM, RANK, IERROR)
```

- Arguments:
 - comm**: communicator (input)
 - rank**: rank of processes (output)

Get_rank in Python

- Rank (label) of the process

```
comm.Get_rank()
```

- No arguments
- Return the rank of this process in a communicator
- (alt. use `comm.rank`)

Copying communicators

- Extensive use of **MPI_COMM_WORLD** is *discouraged*
- Exactly **one** reference to **MPI_COMM_WORLD** in the program (apart from **MPI_Abort**):
- Copy it, e.g.:

```
my_world = MPI_COMM_WORLD
```

- Use **my_world** in the rest of the program
- Declare **my_world** as
 - **MPI_Comm** in C
 - **INTEGER** in Fortran 90
 - **type (MPI_Comm)** in Fortran 08

Copying communicators in Python

- Extensive use of `MPI.COMM_WORLD` is *discouraged*
- Exactly **one** reference to `MPI.COMM_WORLD` in the program (apart from **Abort**):

```
my_world = MPI.COMM_WORLD
```

- Use `my_world` in the rest of the program, example:

```
my_rank = my_world.Get_rank()
```

MPI_Abort

- Aborting all MPI tasks from any task (e.g. read corrupt input file, failed safety check)

```
int MPI_Abort(MPI_Comm comm, int errorcode)
```

In Fortran 90:

```
MPI_ABORT(COMM, ERRORCODE, IERROR)
```

```
INTEGER :: COMM, ERRORCODE, IERROR
```

- **COMM** is the communicator with the task to abort
 - typically `MPI_COMM_WORLD`
- **ERRORCODE** returned to the UNIX shell to flag a problem
 - Return a **1** if you do not understand this
- All arguments: input

Abort in Python

- Aborting all MPI tasks from any task (e.g. read corrupt input file, failed safety check)

```
comm.Abort(errorcode)
```

- Typically called by **MPI.COMM_WORLD**
- **errorcode** returned to the UNIX shell to flag a problem
 - Return a **1** if you do not understand this

Summary

- Concept of communicator
- Predefined communicator `MPI_COMM_WORLD` (`MPI.COMM_WORLD` in Python)
- Querying task rank and size of a communicator
- Aborting a program on error
- You should now be able to write simple MPI programs, which are useful (e.g. task farm)