

Message Passing with MPI

Joachim Hein
LUNARC & Centre of Mathematical Sciences
Lund University

1

Compiling and running MPI code
on COSMOS at LUNARC

2

Overview

- Connecting to COSMOS using Lunarc HPC desktop
- The Lunarc module environment
- Building MPI executables on Lunarc systems
- Executing MPI jobs on Lunarc systems
- Write your first MPI program

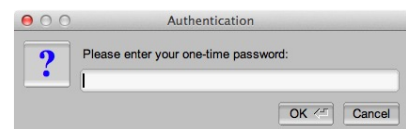
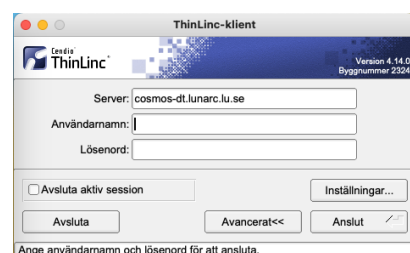
3

Connecting to Aurora: Lunarc HPC Desktop

- Download and install the client from: <http://www.cendio.com/downloads/clients/>
- Launch the client
- Enter "cosmos-dt.lunarc.lu.se" in the server field.
- Enter your user-id & password
- Click [Connect]
- Enter the one time password from your PocketPass app

Detailed descriptions:

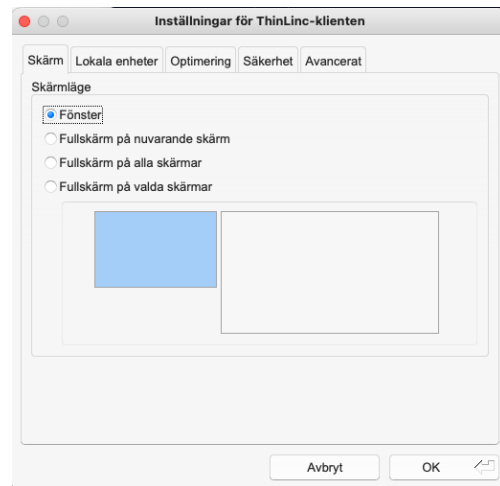
- https://lunarc-documentation.readthedocs.io/en/latest/getting_started/authenticator_howto/
- https://lunarc-documentation.readthedocs.io/en/latest/getting_started/using_hpc_desktop/



4

Some Useful Tricks

- Disable “fullscreen” before connecting
- F8 (fn F8 on MBP) during a DT session could be your friend



5

Building MPI executables

C, C++ and Fortran

6

Building MPI executables

- You use standard compilers (e.g. Intel, GCC, ...) to compile the source
- Link against the MPI library to build an MPI executable
- LUNARC/NAISS systems use modules to make these steps convenient

7

Quick primer to modules

- Show all available modules
`module spider`
- Load a recent foss toolchain (GCC 13.3.0, OpenMPI 5.0.3, ...)
`module load foss/2024a`
- We also provide Intel toolchains
- To get rid of all modules
`module purge`

8

Hands on: Compiling MPI with OpenMPI

- Compile and link your code using wrapper scripts

Language	Command
Fortran	<code>mpifort</code>
C	<code>mpicc</code>
C++	<code>mpiCC</code>

- The wrappers will call the compiler and tell it about the MPI-lib
- For a foss toolchain they will call `gcc`, `g++` or `gfortran`

10

Example: Fortran source with OpenMPI

- Can use options of underlying compiler
- Example:


```
mpifort -O3 -march=native -o program program.f90
```

 - takes the source: `program.f90`
 - calls `gfortran` compiler
 - creates an executable: `program`
 - uses optimisation level: `O3`
 - optimises for present architecture (instruction set)

11

Example: C-source with OpenMPI

- Can use options of underlying compiler
- Example:


```
mpicc -O3 -march=native -o program program.c
```

 - takes the source: `program.f90`
 - calls gcc compiler
 - creates an executable: `program`
 - uses optimisation level: `O3`
 - optimises for present architecture (instruction set)

12

Cmake

- CMake has a utility: **FindMPI**
 - locate the MPI library
 - set up compiler and linker flags
- FindMPI set up the following Cmake variables (v 2.8.12)
 - `MPI_<lang>_FOUND`
 - `MPI_<lang>_COMPILER`
 - `MPI_<lang>_INCLUDE_PATH`
 - `MPI_<lang>_LINK_FLAGS`
 - `MPI_<lang>_COMPILE_FLAGS`
 - `MPI_<lang>_LIBRARIES`
- `<lang>` is one of C, CXX and Fortran

13

CMake example for MPI in Fortran

```
project(mpihello)
enable_language(Fortran)
include(FindMPI)
add_executable(mpihello.x mpihello.f90)

target_include_directories(mpihello.x
                          PUBLIC "${MPI_Fortran_INCLUDE_PATH}")
target_link_libraries(mpihello.x "${MPI_Fortran_LIBRARIES}")
set(CMAKE_Fortran_FLAGS
    "${CMAKE_Fortran_FLAGS} ${MPI_Fortran_COMPILE_FLAGS}")
set(CMAKE_EXE_LINKER_FLAGS
    "${CMAKE_EXE_LINKER_FLAGS} ${MPI_Fortran_LINK_FLAGS}")
```

14

Running MPI programs

- Make sure the modules (compiler and MPI library) used for building the executable are loaded (shared objects!)
- Message passing jobs need to run inside the batch submission system
- Batch submission systems are vital for running message passing applications
- Expect a similar set-up on any well managed service
- Use the course account for quick turn around of jobs

15

Standard MPI job

Run executable: `processor_mpi`

```
#!/bin/bash
#SBATCH -N 2                # number of nodes
#SBATCH --tasks-per-node=48 # number of processes per node
#SBATCH -t 00:05:00        # job-time - here 5 min
#SBATCH -J data_process    # name of job
#SBATCH -o process_mpi_%j.out # output file
#SBATCH -e process_mpi_%j.err # error messages
#SBATCH -A lu2025-7-75      # only for the course
#SBATCH --reservation=mpi-course-1day # only for the course - change for other course days

cat $0
module purge
module load foss/2024a
mpirun --bind-to core ./processor_mpi      #run the program
```

16

Running MPI4PY scripts

17

We need a few extra modules

- Load a recent foss toolchain (GCC 13.3.0, OpenMPI 5.0.3, ...)

```
module load foss/2024a
```

- We need Python, NumPy, SciPy, MPI4PY

```
module load Python/3.12.3
```

```
module load SciPy-bundle/2024.05
```

```
module load mpi4py/4.0.1
```

- To get rid of all modules

```
module purge
```

- Required in jobscript

18

Standard MPI4PY job

Run executable: **processor_mpi.py**

```
#!/bin/bash
#SBATCH -N 2                # number of nodes
#SBATCH --tasks-per-node=48 # number of processes per node
#SBATCH -t 00:05:00         # job-time - here 5 min
#SBATCH -J data_process     # name of job
#SBATCH -o process_mpi_%j.out # output file
#SBATCH -e process_mpi_%j.err # error messages
#SBATCH -A lu2025-7-75      # only for the course
#SBATCH --reservation=mpi-course-1day # only for the course - update as needed

cat $0
module purge
module load foss/2024a
module load Python/3.12.3 SciPy-bundle/2024.05 mpi4py/4.0.1
mpirun --bind-to core python3 ./processor_mpi      #run the python script
```

19

Interacting with SLURM

Submission, queue monitoring and modification

20

Submission with **sbatch**

- Use sbatch to submit your job script to the job-queue
- Example:

```
[fred@cosmos Timetest]$ sbatch runjob.sh  
Submitted batch job 7197
```

- Submit script “runjob.sh”
- Successful submission returns a job-id number

21

Monitoring the queue with **squeue**

- Use **squeue** to monitor the job queue

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
7303	lu48	hybrid_n	fred	PD	0:00	32	(Priority)
7302	lu48	hybrid_n	fred	PD	0:00	32	(Priority)
7301	lu48	hybrid_n	fred	PD	0:00	32	(Resources)
7304	lu48	preproce	karl	PD	0:00	6	(Priority)
7300	lu48	hybrid_n	fred	R	0:24	32	cn[001-032]
7305	lu48	preproce	karl	R	0:37	6	cn[081-086]
7306	lu48	hybrid_n	fred	R	0:37	6	cn[081-086]
7307	lu48	testsimu	sven	R	0:07	1	cn081

- Typically lots of output – use options of **squeue** to filter

22

Options of **squeue**

- Showing jobs for a specific user

```
squeue -u fred
```

will show the jobs of user “fred” only

- Option `--start` gives the estimated job start time
 - This can shift in either direction

23

Deleting jobs with `scancel`

- You can cancel a queued or running job
- Determine job-id, e.g. with `squeue`
- Use `scancel`

```
scancel 7103
```

- terminates job 7103, if running
- removes from the queue

24

Exercise

- See the section “Running and building MPI codes” in the practicalities document
- Make sure you can
 - Connect to COSMOS
 - Compile C and/or Fortran code
 - Interact with SLURM
 - Run MPI code (C, Fortran and/or Python)

25