

# mpi4py and Numpy

Juan de Gracia, Xin Li  
PDC - KTH, Stockholm

# Overview

- Why NumPy with MPI?
- NumPy arrays as buffer-like objects
  - How to check contiguity with flags
  - How to check how memory allocation with strides
- Blocking point-to-point with Send and Recv
- Non-blocking point-to-point with Isend and Irecv
- Collectives I: Bcast
- Collectives II: uneven data with Scatterv and Gatherv
- Collectives III: Reduce

# Introduction to NumPy and MPI

# Introduction to NumPy and MPI

- NumPy:
  - High-performance numerical library.
  - Supports large arrays and vectorized operations.
- MPI:
  - Enables parallel programming for distributed systems.
  - Ideal for scalable computations.
- Together:
  - Communicate and process NumPy arrays efficiently in parallel

# Advantages of Using NumPy with MPI

- **Performance:** Leverages both NumPy's speed and MPI's parallelism.
- **Ease of Use:** Direct integration with `mpi4py`.
- Example Applications:
  - Distributed matrix operations.
  - Large-scale scientific simulations.

# Communication of NumPy Arrays

- **Buffer-Like Objects:** NumPy arrays can directly serve as MPI buffers.
- Some methods:
  - Blocking: `Send`, `Recv`
  - Non-blocking: `Isend`, `Irecv`
  - Collective: `Bcast`, `Scatterv`, `Gatherv`, `Reduce`

# NumPy Arrays as Buffer-Like Objects

# What are Buffer-Like objects

- Arrays used directly as send/receive buffers.
- Efficient and close to MPI's native communication speed.
- Example:
  - `comm.Send([data, MPI.DOUBLE], dest=1)`



# Requirements for Numpy Buffers

- Must be **contiguous** in memory (MPI expects that)
- Receiving buffers must be **pre-allocated** and of the same size as the send buffer
- Example of how to allocate a buffer:
  - `recvbuf = np.zeros(data.size, dtype='float64')`
  - `comm.Recv(recvbuf, source=0)`

# How to check contiguity with .flags

- NumPy flags displays if the array is contiguous in memory and according to which criteria.
- Flags to check:
  - `C_CONTIGUOUS` Row-major (C-Style)
  - `F_CONTIGUOUS` Column major (Fortran-Style)
- By default NumPY uses C-Style
- Try this out:
  - `a = np.random.rand(10,10)`
  - `a.flags`
  - `b = a.T`
  - `b.flags`
  - `c = b.copy()`
  - `c.flags`

# How to check contiguity with .flags


- NumPy flags displays if the array is contiguous in memory and according to which criteria.
- Flags to check:
  - `C_CONTIGUOUS` Row-major (C-Style)
  - `F_CONTIGUOUS` Column major (Fortran-Style)
- By default NumPY uses C-Style
- Try this out:
  - `a = np.random.rand(10,10)`
  - `a.flags # C_CONTIGUOUS = True`
  - `b = a.T`
  - `b.flags # F_CONTIGUOUS = True`
  - `c = b.copy()`
  - `c.flags # C_CONTIGUOUS = True`

# How to check memory allocation with Strides


- What are strides?
  - Indicates how many **bytes** to move in each dimension to access the next element.
  - The method `a.strides` will return a tuple with:
    - For 1D arrays the size of the data type in bytes (1,2,4,8 = 8,16,32,64)
    - For 2D arrays (row stride, column stride)
      - C Style (`row_stride > column_stride`)
      - F Style (`row_stride < column_stride`)

# Understanding C-Style vs F-Style

**C Style = Row-major**



**Fortran-Style = Column major**



# Use NumPy array as buffer-like object

- The array itself, or a list or tuple with
  - 2 or 3 elements
  - 4 elements for the vector variants (Scatterv, Gatherv)
- data
- [data, MPI.DOUBLE]
- [data, n, MPI.DOUBLE]
- [data, count, displ, MPI.DOUBLE]

Blocking send/recv

# Send / Recv

- Syntax
  - `comm.Send(obj, dest=rank, tag=tag)`
  - `comm.Recv(obj, source=rank, tag=tag)`
- Note
  - obj needs to be created prior to the communication
  - size of obj needs to be known before hand



# Send / Recv

- example

```
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

if rank == 0:
    data = np.arange(10, dtype=np.float64)
    for i in range(1, size):
        comm.Send(data, dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(10, dtype=np.float64)
    comm.Recv(data, source=0, tag=rank)
    print(f"Rank {rank} received data from rank 0: {data}")

print(f"Rank {rank} has data: {data}")
```

# Send / Recv

- output

Rank 0 sent data to rank 1

Rank 0 sent data to rank 2

Rank 0 sent data to rank 3

Rank 0 has data: [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]

Rank 1 received data from rank 0: [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]

Rank 1 has data: [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]

Rank 2 received data from rank 0: [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]

Rank 2 has data: [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]

Rank 3 received data from rank 0: [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]

Rank 3 has data: [0. 1. 2. 3. 4. 5. 6. 7. 8. 9.]

# Send / Recv

- exercise
  - what will happen if you try to send an array with non-contiguous memory?
  - hint: this is how to create a simple array with non-contiguous memory  
`data = np.arange(12.)[:, :2]`

# Send / Recv

- exercise
  - what will happen if the receiving buffer is *larger* than the sent array?

```
if rank == 0:
    data = np.arange(10, dtype=np.float64)
    for i in range(1, size):
        comm.Send(data, dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(12, dtype=np.float64)
    comm.Recv(data, source=0, tag=rank)
    print(f"Rank {rank} received data from rank 0: {data}")

print(f"Rank {rank} has data: {data}")
```

# Send / Recv

- exercise
  - what will happen if the receiving buffer is *smaller* than the sent array?

```
if rank == 0:
    data = np.arange(10, dtype=np.float64)
    for i in range(1, size):
        comm.Send(data, dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(9, dtype=np.float64)
    comm.Recv(data, source=0, tag=rank)
    print(f"Rank {rank} received data from rank 0: {data}")
```

# Send / Recv

- best practice: check status

```
if rank == 0:
    data = np.arange(10, dtype=np.float64)
    for i in range(1, size):
        comm.Send(data, dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(9, dtype=np.float64)
    status = MPI.Status()
    comm.Recv(data, source=0, tag=rank, status=status)
    if status.error != 0:
        comm.Abort(status.error)

    print(f"Rank {rank} received data from rank 0: {data}")

print(f"Rank {rank} has data: {data}")
```

# Send / Recv with buffer size

- use [data, n, MPI.DOUBLE] to specify the buffer

```
if rank == 0:
    data = np.arange(10, dtype=np.float64)
    for i in range(1, size):
        # Using [data, size, type] to send a buffer of data
        comm.Send([data, 2, MPI.DOUBLE], dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(10, dtype=np.float64)
    # Using [data, size, type] to receive a buffer of data
    comm.Recv([data, 2, MPI.DOUBLE], source=0, tag=rank)

    print(f"Rank {rank} received data from rank 0: {data[:2]}")

print(f"Rank {rank} has data: {data}")
```

# Reminder: Standard send in C: MPI\_Send

```
int MPI_Send(void* buf, int count, MPI_Datatype  
datatype, int dest, int tag, MPI_Comm comm)
```

- `buf`: address of send buffer
- `count`: number of elements to be sent
- `datatype`: data type of buffer (explained further down)
- `dest`: rank of receiver
- `tag`: message tag (put 0 if you don't need)
- `comm`: communicator



# Send / Recv with buffer size

- use array slicing

```
if rank == 0:
    data = np.arange(10, dtype=np.float64)
    for i in range(1, size):
        # Using [data, size, type] to send a buffer of data
        comm.Send([data[2:5], 2, MPI.DOUBLE], dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(10, dtype=np.float64)
    # Using [data, size, type] to receive a buffer of data
    comm.Recv([data[2:5], 2, MPI.DOUBLE], source=0, tag=rank)

    print(f"Rank {rank} received data from rank 0: {data}")

print(f"Rank {rank} has data: {data}")
```

# Send / Recv with 2D array

- exercise
  - What will happen if we send a 2D array with .T (transpose)?

```
if rank == 0:
    data = np.arange(16, dtype=np.float64).reshape(4, 4)
    for i in range(1, size):
        comm.Send(data, dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(16, dtype=np.float64).reshape(4, 4)
    status = MPI.Status()
    comm.Recv(data, source=0, tag=rank)

    print(f"Rank {rank} received data from rank 0: {data}")

print(f"Rank {rank} has data: {data}")
```

# Send / Recv with 2D array

- exercise
  - What will happen if we send a 2D array with `.T.copy()` (copy of transpose)?

```
if rank == 0:
    data = np.arange(16, dtype=np.float64).reshape(4, 4).T.copy()
    for i in range(1, size):
        comm.Send(data, dest=i, tag=i)
        print(f"Rank 0 sent data to rank {i}")

else:
    data = np.empty(16, dtype=np.float64).reshape(4, 4)
    status = MPI.Status()
    comm.Recv(data, source=0, tag=rank)

    print(f"Rank {rank} received data from rank 0: {data}")

print(f"Rank {rank} has data: {data}")
```

Non-blocking send/recv

# Isend / Irecv

- Syntax
  - `comm.Isend(obj, dest=dest, tag=tag)`
  - `comm.Irecv(obj, source=source, tag=rag)`
- Note
  - obj needs to be created prior to the communication
  - size of obj needs to be known before hand
  - A Request object is returned by Isend / Irecv
  - Use Wait method of the Request object
  - No Status involved

# Isend / Irecv

- example

```
# Example non blocking Isend/Irecv
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

if rank == 0:
    data = np.arange(10, dtype=np.float64)
    reqs = []
    for i in range(1, size):
        reqs.append(comm.Isend(data, dest=i, tag=i))
    for req in reqs:
        req.Wait()
    print('process 0 sent:', data)
else:
    data = np.empty(10, dtype=np.float64)
    req = comm.Irecv(data, source=0, tag=rank)
    req.Wait()
    print('process', rank, 'received:', data)
```

# Isend / Irecv

- exercise:
  - what if the size of the receiving buffer is larger than the sent array?
  - what if the size of the receiving buffer is smaller than the sent array?
- send a 2D array
  - without transpose
  - with .T
  - with .T.copy()

# Collectives



# Bcast

- Syntax
  - `comm.Bcast(obj, root=root)`
- Note
  - obj needs to be created prior to the communication
  - size of obj needs to be known before hand
  - root can be non-zero (source of communication)

# Bcast

# Example of Bcast

```
import numpy as np
from mpi4py import MPI

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

if rank == 0:
    data = np.arange(10, dtype='i')
else:
    data = np.empty(10, dtype='i')

comm.Bcast(data, root=0)

print("Rank: ", rank, "has data: ", data)
```

# Bcast

- exercise
  - what will happen if the size of the receiving buffer is *larger*?
  - what will happen if the size of the receiving buffer is *smaller*?

# Scatterv

- vector variant of Scatter
- Syntax
  - `comm.Scatterv([sendbuf, count, displ, MPI.DOUBLE], recvbuf, root=root)`
- Note
  - `[sendbuf, count, displ, MPI.DOUBLE]` defines the sending buffer
  - `recvbuf` needs to be created prior to the communication
  - size of `recvbuf` needs to be known before hand
  - `root` can be non-zero (source of Scatterv)

# Scatterv

```
# Example of Scatterv and Gatherv
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

if rank == 0:
    sendbuf = np.arange(100, dtype='i')
    # Determine the counts and displacements
    # for the scatter operation
    counts = np.full(size, 100//size, dtype='i')
    counts[:100%size] += 1
    displs = np.insert(np.cumsum(counts), 0, 0)[:size]

    print('counts:', counts)
    print('displs:', displs)

else:
    sendbuf = None
    counts = np.empty(size, dtype='i')
    displs = None

# Scatter the data
# First we broadcast the counts to all processes
comm.Bcast(counts, root=0)

# Otherwise the recvbuf will have the same size as the sendbuf
recvbuf = np.empty(counts[rank], dtype='i')

comm.Scatterv([sendbuf, counts, displs, MPI.INT], recvbuf, root=0)

print('Rank:', rank, 'recvbuf:', recvbuf)
```

# Scatterv

- output

```
counts: [25 25 25 25]
```

```
displs: [ 0 25 50 75]
```

```
Rank: 0 recvbuf: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24]
```

```
Rank: 1 recvbuf: [25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49]
```

```
Rank: 2 recvbuf: [50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74]
```

```
Rank: 3 recvbuf: [75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99]
```

# Gatherv

- vector variant of Gather
- Syntax
  - `comm.Gatherv(sendbuf, [recvbuf, count, displ, MPI.DOUBLE], root=root)`
- Note
  - `[recvbuf, count, displ, MPI.DOUBLE]` defines the receiving buffer
  - `recvbuf` needs to be created prior to the communication
  - size of `recvbuf` needs to be known before hand
  - `root` can be non-zero (target of Gatherv)

# Gatherv

```
# Example of Scatterv and Gatherv
```

```
from mpi4py import MPI
import numpy as np
```

```
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()
```

```
if rank == 0:
    sendbuf = np.arange(100, dtype='i')
    counts = np.full(size, 100 // size, dtype='i')
    counts[:100 % size] += 1
    displs = np.zeros(size, dtype='i')
    displs[1:] = np.cumsum(counts[:-1])

    print('counts:', counts)
    print('displs:', displs)
else:
    sendbuf = None
    counts = np.empty(size, dtype='i')
    displs = np.empty(size, dtype='i')
```

```
# Broadcast counts and displacements to all ranks
```

```
comm.Bcast(counts, root=0)
comm.Bcast(displs, root=0)
```

```
# Allocate recvbuf based on counts for the current rank
```

```
recvbuf = np.empty(counts[rank], dtype='i')
```

```
# Scatter the data
```

```
comm.Scatterv([sendbuf, counts, displs, MPI.INT], recvbuf, root=0)
print(f'Rank {rank} received data: {recvbuf}')
```

```
# Gather the data back
```

```
sendbuf2 = recvbuf
```

```
if rank == 0:
```

```
    recvbuf2 = np.empty(100, dtype='i') # Full array to gather data
```

```
else:
```

```
    recvbuf2 = None
```

```
comm.Gatherv(sendbuf2, [recvbuf2, counts, displs, MPI.INT], root=0)
```

```
if rank == 0:
```

```
    print('Gathered data:', recvbuf2)
```



# Gatherv

- output

```
Gathered data: [ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14
15 16 17 18 19 20 21 22 23
24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43
44 45 46 47
48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67
68 69 70 71
72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91
92 93 94 95
96 97 98 99]
```

# Reduce

- Syntax
  - `comm.Reduce(sendbuf, recvbuf, op=op, root=root)`
- Note
  - default op is `MPI.SUM`
  - root can be non-zero (target of Reduce)

# Reduce

# Example of Scatterv and Reduce

```
from mpi4py import MPI
import numpy as np
```

```
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()
```

```
if rank == 0:
    sendbuf = np.arange(100, dtype='i')
    # Determine counts and displacements
    counts = np.full(size, 100 // size, dtype='i')
    counts[:100 % size] += 1
    displs = np.zeros(size, dtype='i')
    displs[1:] = np.cumsum(counts[:-1])
```

```
    print('counts:', counts)
    print('displs:', displs)
```

```
else:
```

```
    sendbuf = None
    counts = np.empty(size, dtype='i')
    displs = np.empty(size, dtype='i')
```

```
# Broadcast counts and displacements to all ranks
```

```
comm.Bcast(counts, root=0)
comm.Bcast(displs, root=0)
```

```
# Allocate recvbuf based on counts for the current rank
recvbuf = np.empty(counts[rank], dtype='i')
```

```
# Scatter the data
```

```
comm.Scatterv([sendbuf, counts, displs, MPI.INT], recvbuf, root=0)
```

```
print(f'Rank {rank} received data: {recvbuf}')
```

```
# Compute the partial sum on each rank
```

```
partial_sum = np.sum(recvbuf) # This is a scalar value
```

```
# Use Reduce to compute the total sum across all ranks
```

```
total_sum = comm.reduce(partial_sum, op=MPI.SUM, root=0)
```

```
if rank == 0:
```

```
    print('Total sum of all elements:', total_sum)
```

# Reduce

- output

Rank 0 received data: [ 0 1 2 3 4 5 6 7 8 9  
10 11 12 13 14 15 16 17 18 19 20 21 22 23  
24]

Rank 1 received data: [25 26 27 28 29 30 31 32 33 34  
35 36 37 38 39 40 41 42 43 44 45 46 47 48  
49]

Rank 2 received data: [50 51 52 53 54 55 56 57 58 59  
60 61 62 63 64 65 66 67 68 69 70 71 72 73  
74]

Rank 3 received data: [75 76 77 78 79 80 81 82 83 84  
85 86 87 88 89 90 91 92 93 94 95 96 97 98  
99]

Total sum of all elements: 4950

# Reduce

- exercise
  - Use Reduce to compute the sum of numpy arrays

# Reduce

```
# Simple example of Reduce
from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

data = np.array([1, 2, 3, 4], dtype=np.float64) * (rank + 1)
result = np.empty(4, dtype=np.float64)

# Print partial data
print(f"Rank {rank} has data: {data}")

comm.Reduce(data, result, op=MPI.SUM, root=0)

if rank == 0:
    print(f"Rank {rank} has data after reduce: {result}")
```

# Reduce

Rank 1 has data: [2. 4. 6. 8.]

Rank 2 has data: [ 3. 6. 9. 12.]

Rank 3 has data: [ 4. 8. 12. 16.]

Rank 0 has data: [1. 2. 3. 4.]

Rank 0 has data after reduce: [10. 20. 30. 40.]

# Summary



# Numpy array as buffer-like object in mpi4py

- fast communication
- less flexible code (need to deal with memory)
- blocking and nonblocking communication
  - Send, Recv, Isend, Irecv
- collective communication
  - Bcast, Scatterv, Gatherv, Reduce